


I deployed my fine-tuned ticket router on Cloud Run

TicketRouter is a Qwen3-1.7B + LoRA model served through a containerized FastAPI endpoint.

A screenshot of a web browser window displaying the response from the TicketRouter API endpoint. The browser's address bar shows the URL 'ticketrouter-cpu-api-281584520321.us-east4.run.app'. Below the address bar, there is a 'Pretty-print' checkbox which is checked. The main content area shows a JSON response with the following structure: { "service": "TicketRouter API", "status": "ok", "model_id": "Neog007/TicketRouter-1.7B", "model_loaded": false, "cuda_available": false, "endpoints": { "health": "/health", "route": "/route" }, "example": { "method": "POST", "path": "/route", "body": { "ticket": "My Outlook calendar invites are not showing up in Teams meetings." } } }.

```
{
  "service": "TicketRouter API",
  "status": "ok",
  "model_id": "Neog007/TicketRouter-1.7B",
  "model_loaded": false,
  "cuda_available": false,
  "endpoints": {
    "health": "/health",
    "route": "/route"
  },
  "example": {
    "method": "POST",
    "path": "/route",
    "body": {
      "ticket": "My Outlook calendar invites are not showing up in Teams meetings."
    }
  }
}
```

Live API

Public Cloud Run endpoint with service metadata

The deployment path mattered as much as the model score.

The goal was not to stop at a notebook. I wanted a small production-shaped service that could accept tickets and return routing decisions.

1

Qwen3-1.7B + LoRA

Run 2 model artifact

2

FastAPI

health and route endpoints

3

Docker

repeatable container image

4

Artifact Registry

versioned image storage


5

Cloud Run


public serverless API

The API exposes the exact proof points I need for a demo.

A simple service contract keeps the demo honest: status, health, and inference.



GET /health



POST /route

The route endpoint sends an Outlook and Teams ticket to the model and returns O365 with confidence and an auto-route action.

The deployed service is backed by controlled fine-tuning results.

Baseline


24.8%

Run 2

77.8%


Macro F1

0.753



Run 2 won across four experiments and became the production candidate.

What made it production-minded: constraints, fallback, and honesty.



Deployment notes

CPU Cloud Run demo is live and public.

NVIDIA L4 GPU config was prepared, but quota limited execution.

Low-confidence tickets route to human review.

Active Directory remains the weakest class and needs better labeled data.

Takeaway

A useful AI project is not just a model. It is a service with clear evaluation, known failure modes, and a deployment path.